

CS5361/6361 Machine Learning

Fall 2024

Arrays – Exercise 1

Write the following functions. See the commented out part of the program for expected results. Also, make sure your functions don't modify the array (if any) received as parameter.

1. Write the function `is_square(X)` that receives an array `X` and determines if `X` is square. An array is square if it has two dimensions and they are the same size (that is, it has the same number of rows and columns).
2. The numpy function `np.max(X)` returns the maximum value in array `X`. Write the function `max_loop(X)` that uses a loop to find the maximum value in array `X`, where `X` may have any number of dimensions.
3. The numpy function `np.argmax(X)` returns the index where the maximum value in array `X` is stored. Write the function `argmax_loop(X)` that uses a loop to find the argmax in 1D array `X`.
4. The numpy function `np.sum(X)` returns the sum of the elements in array `X`. Write the function `sum_loop(X)` that uses a loop to compute the sum of the elements in array `X`, where `X` may have any number of dimensions.
5. Write the function `diagonal(X)` that receives a square array `X` and returns a 1D array containing the elements in the diagonal of `X` (that is `[X[0,0], X[1,1], ...]`).
6. Write the function `count_digits(X)` that counts the number of times each of the numbers 0,1,...,9 appears in an array. The function should receive an array of integers `X` of any dimensionality and return a 1D array of length 10, where the first element in the array is the number of times 0 appears in `X`, the second element is the number of times 1 appears in `X`, and so on.
7. The numpy function `np.dot(x1,x2)` receives two 1D arrays `x1` and `x2` and returns their dot product. Write the function `dot(x1,x2)` that computes the dot product of `x1` and `x2` without using the numpy function.
8. Write the function `euclidean_dist(x1,x2)` that receives 1D arrays `x1` and `x2` and returns the Euclidean distance from the point represented by `x1` to the point represented by `x2`.
9. Write the function `manhattan_dist(x1,x2)` that receives 1D arrays `x1` and `x2` and returns the Manhattan Euclidean distance from the point represented by `x1` to the point represented by `x2`.
10. Write the function `accuracy(p,y)` that receives 1D arrays of integers `p` and `y`, where `(p.shape == y.shape)` and returns the accuracy of a classifier whose predictions are given by `p` and where the correct classification is given by `y`.
11. Write the function `mse(p,y)` that receives 1D float arrays `p` and `y`, where `(p.shape == y.shape)` and returns the mean-squared error of a model whose predictions are given by `p` and where the correct target value is given by `y`.
12. Write the function `select_features(X,n)` that receives 2D array `X` representing a dataset and an integer `n` and returns a subset of `X` that contains all instances in `X` but only the features with highest variance (you may use the `np.var` numpy function).
13. Write the function `select_instances(X,f,t)` that receives 2D array `X` representing a dataset, an integer `f` representing a feature (or column number) in `X` and a floating point value `t` and returns a subset of `X` containing only the instances where feature `f` is larger than `t`.

14. Write the function `nearest_neighbor(X,x)` that receives a 2D array `X` of shape (n,c) and a 1D array `x` of shape $(c,)$ and returns the index of the row in `X` that is most similar (according to Euclidean distance) to `x`.
15. Write the function `nearest_neighbors(X1,X2)` that receives 2D arrays `X1` and `X2`, where $(X1.shape[1] == X2.shape[1])$ and returns a 1D array `N` of shape $X2.shape[0]$, where `N[i]` is the instance in `X1` that is most similar to `X2[i]` according to Euclidean distance. Do this using at most one loop; for extra credit do it with no loops.